

ARTIFICIAL INTELLIGENCE

Complete Final Exam Study Guide

State Space Search · Heuristic Search · Adversarial Search
Local & Evolutionary Search · Machine Learning · ANN & Backpropagation
CNN · Clustering · Performance Evaluation

Extreme Detail · TikZ Diagrams & PGFPlots · Dry Runs · Easy to Hard Exam Questions

Backpropagation: Sigmoid Only · No Python, Pure LaTeX


0. 1

1	State Space Search – Foundations	3
2	Blind / Uninformed Search	3
2.1	Depth First Search (DFS)	3
2.1.1	DFS Full Tree Diagram with Traversal Order	4
2.2	Breadth First Search (BFS)	4
2.2.1	BFS Level-by-Level Diagram	5
2.2.2	BFS vs DFS Memory Usage – Graph	5
2.3	Iterative Deepening Search (IDS)	6
2.3.1	IDS Iteration Visualization	6
2.3.2	IDS vs BFS: Re-expansion Cost and Space Saving	6
2.4	Uniform Cost Search (UCS)	7
2.4.1	UCS Weighted Graph Dry Run	7
2.5	Bidirectional Search	7
2.6	All Blind Search Algorithms Compared	8
3	Heuristic / Informed Search	8
3.1	Greedy Best-First Search	8
3.2	A* Search	8
3.2.1	Romania Map: Greedy vs A* Side-by-Side	8
3.3	Admissibility and Consistency	9
3.4	Heuristic Quality: Dominance	9
4	Adversarial Search (Game Playing)	10
4.1	Minimax Algorithm	10
4.1.1	Minimax Full Tree Dry Run	10
4.2	Alpha-Beta Pruning	10
4.2.1	Alpha-Beta Full Dry Run with Three MAX Children	10
4.2.2	Alpha-Beta Best/Worst/Average Case Graph	11
5	Local Search – Hill Climbing	11
5.0.1	Objective Function Landscape – Problems Illustrated	12
5.0.2	Steepest Ascent HC Dry Run with Graph	12
6	Evolutionary / Genetic Algorithms	13
6.0.1	GA Cycle Diagram	13
6.1	Selection Mechanisms	13
6.1.1	Roulette Wheel (Fitness Proportionate) Selection	14
6.1.2	Tournament Selection	14
6.1.3	Rank-Based Selection	14
6.2	Genetic Operators	14
6.2.1	Single-Point Crossover	14
6.2.2	Bit-Flip Mutation	14
6.2.3	GA Fitness Over Generations Graph	15
7	Machine Learning and Linear Regression	15
7.0.1	Linear Regression Dry Run with Scatter Plot	15
8	Artificial Neural Networks (ANN)	16
8.1	Network Architecture	16
8.2	Activation Functions	16
8.2.1	All Major Activation Functions	16
8.3	Perceptron and Linear Separability	17
9	Backpropagation (Sigmoid Only)	17
9.1	Complete Backprop Dry Run – Full Numerical Example	18
9.1.1	Network Architecture Diagram	18
9.1.2	Loss Reduction Over Training Epochs	20
9.2	Multi-Hidden-Layer Generalization	20

10 ANN Training Issues	20
10.1 Overfitting, Underfitting, and Bias-Variance Tradeoff	21
10.2 Vanishing Gradient Problem with Sigmoid	21
10.3 Dropout Regularization	22
11 Performance Evaluation	22
11.1 Confusion Matrix	22
11.1.1 Numeric Dry Run	22
11.1.2 Precision vs Recall Trade-off	23
11.2 Multi-Class Metrics: Macro and Micro Averaging	23
12 Convolutional Neural Networks (CNN)	23
12.1 Convolution Operation Dry Run	24
12.2 Max Pooling Dry Run	24
12.2.1 Full CNN Architecture Sizes and Parameters	24
13 Clustering	25
13.1 K-Means Clustering – Complete Dry Run	25
13.2 Agglomerative (Hierarchical) Clustering	26
13.2.1 Complete-Linkage Agglomerative Dry Run	26

1. 1

A **state space** is a 5-tuple: $\langle S, s_0, A, T, G \rangle$ where S =all states, s_0 =initial state, A =actions/operators, T =transition function, G =goal states. A **path** is a sequence of states; a **solution** is any path from s_0 to $g \in G$; an **optimal solution** has minimum cost.



Paths: $S \rightarrow G = 8$ — $S \rightarrow A \rightarrow B \rightarrow G = 6$ (optimal) — $S \rightarrow A \rightarrow C \rightarrow G = 11$

Nodes=states, edges=operators with costs. Optimal path is $S \rightarrow A \rightarrow B \rightarrow G$ with total cost 6, not the direct edge costing 8.

Four Key Properties of Any Search Algorithm

Property	Definition
Completeness	Guarantees finding a solution <i>if one exists</i>
Optimality	Guarantees the solution has <i>minimum cost</i>
Time Complexity	Nodes expanded as function of b (branching factor), d (shallowest goal depth), m (max depth)
Space Complexity	Maximum nodes held in memory simultaneously

Key Terminology

Frontier (Open List): Nodes discovered but not yet expanded.

Explored (Closed List): Nodes already expanded (graph search only).

Branching factor b : Average number of successors per node.

Depth d : Number of edges from root to shallowest goal.

Max depth m : Maximum depth of any path in the search tree.

2. 1

Blind search has **no information** about goal location. It only knows whether a state is a goal or not.

2.1 1

DFS explores as **deep as possible** before backtracking. Uses a **stack (LIFO)**.

Key Insight: When you push successors and immediately pop, you always go deeper. The stack only holds the current path plus unexplored siblings at each level — hence linear space $O(bm)$.


DFS Algorithm

1. Push initial state onto stack
2. If stack empty $\Rightarrow$ return failure
3. Pop top node n
4. If n is goal $\Rightarrow$ return solution
5. Mark n visited; push unvisited successors
6. Go to step 2

DFS Properties

Property	Value
Complete?	No (infinite loops)
Optimal?	No
Time	$O(b^m)$
Space	$O(bm)$ – best!

2.1.1 1DFS Full Tree Diagram with Traversal Order



Complete DFS Trace: Initial

Step	Stack (top first)	Visited	Action
0	[A]	{}	Initialize: push A
1	[C, B]	{A}	Pop A; not goal; push successors C,B (B on top)
2	[E, D, C]	{A,B}	Pop B; not goal; push D,E
3	[D, C]	{A,B,E}	Pop E; leaf, not goal
4	[C]	{A,B,E,D}	Pop D; leaf, not goal
5	[G, F]	{A,B,E,D,C}	Pop C; not goal; push F,G
6	[G]	{A,B,E,D,C,F}	Pop F = GOAL! Return path: A→C→F

Note on push

order: We push right child first so left child is on top and explored first. Here we push B before C so B is popped first — deeper left branch explored before right.

Q: [Easy] What data structure does DFS use? Why does it explore depth-first?

A: A Stack (LIFO). Successors are immediately on top, so DFS dives deeper before trying siblings.

Q: [Medium] DFS tree: $b = 3$, max depth $m = 4$. How many nodes in memory at once?

A: $O(bm) = 3 \times 4 = 12$ nodes. At depth k , stack holds: current path ($k + 1$ nodes) + at most $b - 1$ siblings per level $= k + 1 + (b - 1)k = bk + 1$. For $b = 3$, $m = 4$: $3 \times 4 + 1 = 13$ maximum. DFS's **linear space** is its greatest strength.

Q: [Hard] Graph: $1 \rightarrow 2,3 \rightarrow 2 \rightarrow 4,5 \rightarrow 3 \rightarrow 6,7$ (unit costs). Goal=6. Trace DFS left-first. Is path optimal? Why not guaranteed?

A: Stack=[1]. Pop 1 $\rightarrow$ push [3,2]. Pop 2 $\rightarrow$ push [5,4,3]. Pop 4=leaf. Pop 5=leaf. Pop 3 $\rightarrow$ push [7,6]. Pop 6=GOAL. Path: $1 \rightarrow 3 \rightarrow 6$, depth=2, cost=2. Happens to be optimal here. But DFS gives NO general optimality guarantee — it could find $1 \rightarrow 3 \rightarrow 7$ first if push order were different, or a deeper path first in a weighted graph.

2.2 1

BFS explores all nodes at depth d before depth $d + 1$. Uses a **queue (FIFO)**.


Key Insight: Since BFS explores level by level, the first time it reaches the goal it must have taken the fewest edges — hence optimal for uniform costs.

BFS Algorithm

1. Enqueue initial state
2. If queue empty $\Rightarrow$ return failure
3. Dequeue front node n
4. If n is goal $\Rightarrow$ return solution
5. Enqueue all unvisited successors (L $\rightarrow$ R)
6. Go to step 2


BFS Properties

Property	Value
Complete?	Yes
Optimal?	Yes (equal costs)
Time	$O(b^d)$
Space	$O(b^d)$ – worst!

2.2.1 1BFS Level-by-Level Diagram**BFS Trace: same graph**

Step	Queue (front first)	Visited	Action
0	[A]	{}	Enqueue A
1	[B, C]	{A}	Dequeue A $\rightarrow$ enqueue B,C
2	[C, D, E]	{A,B}	Dequeue B $\rightarrow$ enqueue D,E
3	[D, E, F, G]	{A,B,C}	Dequeue C $\rightarrow$ enqueue F,G
4	[E, F, G]	{A,B,C,D}	Dequeue D $\rightarrow$ leaf
5	[F, G]	{A,B,C,D,E}	Dequeue E $\rightarrow$ leaf
6	[G]	{...,F}	Dequeue F = GOAL! Path: A$\rightarrow$C$\rightarrow$F

BFS dequeued 6 nodes; DFS found the goal in fewer steps (by luck of path order). But BFS **guarantees the shallowest** (fewest-edge) path. DFS offers no such guarantee.

2.2.2 1BFS vs DFS Memory Usage – Graph

Q: [Easy] Why is BFS optimal for equal-cost but NOT for weighted graphs?

A: BFS minimizes the *number of edges* (steps). In weighted graphs, fewer edges $\neq$ lower total cost. Example: A $\rightarrow$ B cost=100, A $\rightarrow$ C $\rightarrow$ B cost=1+1=2. BFS finds A $\rightarrow$ B first (1 edge) but it is not optimal.

Q: [Medium] $b = 4$, goal at depth $d = 5$. Worst-case BFS nodes in memory?

A: $O(b^d) = 4^5 = 1024$ nodes. All nodes at depth d must be in the queue simultaneously, plus all at $d - 1$ that haven't been expanded yet.

Q: [Hard] Graph: $1 \rightarrow 2, 3, 4 \rightarrow 5, 6 \rightarrow 7 \rightarrow 8$. Goal=7. Compare BFS vs DFS (left-first). Which expands fewer nodes? Which guarantees optimality?

A: BFS: enqueue[2,3,4], dequeue 2→enqueue[3,4,5,6], dequeue 3→enqueue[4,5,6,7], dequeue 4, dequeue 5, dequeue 6, dequeue 7=GOAL. 7 dequeues. **DFS** (left-first push: right child first): Stack=[1]. Pop 1→push[4,3,2]. Pop 2→push[6,5,4,3]. Pop 5=leaf. Pop 6=leaf. Pop 4→push[8,4,3]. Pop 8=leaf. Pop 4=visited skip. Pop 3→push[7,4]. Pop 7=GOAL. DFS expands fewer nodes here. But BFS guarantees shortest path $1 \rightarrow 3 \rightarrow 7$ (depth 2). DFS has no such guarantee.

2.3 1

IDS runs DFS repeatedly with increasing depth limits $L = 0, 1, 2, \dots$ until goal found. Gets the **best of both**: DFS space + BFS completeness and optimality.


2.3.1 1IDS Iteration Visualization

L=0: Expand only root A. Not found. (1 node)

L=1: Expand A; push B,C. Pop B (dead end depth 1); pop C (dead end). Not found. (3 nodes)

L=2: A→B→D(dead), B→E(dead); C→F=GOAL! Path: A→C→F (6 nodes explored)

2.3.2 1IDS vs BFS: Re-expansion Cost and Space Saving



IDS expands $\frac{b}{b-1}$ times more nodes than BFS. For $b = 3$: 33% overhead, $b = 10$: 11% overhead. Re-expansion penalty is tiny compared to the massive space saving.

IDS Node Count Formula: b

Each node at depth k is re-expanded $(d - k + 1)$ times across all iterations.

$$\text{Total} = (d+1) \cdot b^0 + d \cdot b^1 + (d-1) \cdot b^2 + \dots + 1 \cdot b^d \\ = 5(1) + 4(3) + 3(9) + 2(27) + 1(81) = 5 + 12 + 27 + 54 + 81 = \mathbf{179}$$

BFS would be $1 + 3 + 9 + 27 + 81 = 121$. IDS uses 48% more expansions but uses only $O(bd) = 12$ memory vs $O(b^d) = 81$.

Q: [Easy] What two algorithms does IDS combine? What does it inherit from each?

A: DFS (linear space $O(bm)$) + BFS (completeness + optimality for uniform costs).

Q: [Medium] IDS $b = 3$, $d = 4$. Total nodes generated? Compare to BFS.

A: IDS: 179. BFS: 121. IDS generates 48% more nodes but uses only $O(bd) = 12$ memory vs BFS's $O(b^d) = 81$.

Q: [Hard] When does IDS fail to find the optimal solution?

A: Non-uniform step costs. Example: A→B cost=1, A→C cost=1, B→G cost=100, C→G cost=1. Both

goals at depth 2. IDS finds B→G (cost 101) before seeing C→G (cost 2) if B is explored first. Use UCS for weighted graphs.

2.4 1

UCS always expands the node with **lowest cumulative path cost** $g(n)$. Uses a **priority queue (min-heap)** ordered by $g(n)$. Generalizes BFS to weighted graphs.

Critical rule: Only declare goal when **dequeued**, not when inserted into the priority queue.

UCS Algorithm

1. Insert start with $g(s_0) = 0$ into PQ

2. If PQ empty $\Rightarrow$ failure

3. Extract node n with min g

4. If n is goal $\Rightarrow$ return $g(n)$

5. For each successor s : $g(s) = g(n) + cost(n, s)$; insert/update in PQ


6. Go to step 2

UCS Properties

Property	Value
Complete?	Yes (costs>0)
Optimal?	Yes
Time	$O(b^{1+\lceil C^*/\epsilon \rceil})$
Space	$O(b^{1+\lceil C^*/\epsilon \rceil})$

C^* =optimal cost, ϵ =min step cost.

2.4.1 1UCS Weighted Graph Dry Run




UCS Step-by-Step Trace

Step	Priority Queue (node:g)	Expanded	Notes
0	[(A:0)]	–	Initialize
1	[(C:2),(B:4)]	A($g=0$)	C=0+2=2, B=0+4=4
2	[(B:3),(D:10)]	C($g=2$)	B via C=2+1=3<4: update. D=2+8=10
3	[(D:6),(C:8),(D:10)]	B($g=3$)	D via B=3+3=6, C via B=3+5=8
4	[rest...]	D($g=6$)	GOAL! $g=6$. Optimal path: A→C→B→D

son: A→B→D=4+3=7. A→C→D=2+8=10. A→C→B→D=2+1+3=6. UCS correctly found the optimal.

Compari-

2.5 1



2.6 1

Algorithm	Data Structure	Complete?	Optimal?	Time	Space
DFS	Stack (LIFO)	No*	No	$O(b^m)$	$O(bm)$
BFS	Queue (FIFO)	Yes	Yes**	$O(b^d)$	$O(b^d)$
IDS	Stack (repeated)	Yes	Yes**	$O(b^d)$	$O(bd)$
UCS	PQ by $g(n)$	Yes	Yes	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Bidir.	Two frontiers	Yes***	Yes***	$O(b^{d/2})$	$O(b^{d/2})$

*DFS complete on finite acyclic graphs with cycle detection. **Optimal for uniform step costs only. ***Under certain conditions (known goal, reversible actions).

3. 1

Informed search uses a **heuristic function** $h(n)$: an estimate of cost from n to the nearest goal. $h(\text{goal}) = 0$. A good heuristic can dramatically reduce nodes explored from exponential to near-linear.

3.1 1

$f(n) = h(n)$ (ignores actual cost $g(n)$ already incurred)

Greedy always expands the node that *looks* closest to the goal. It is fast but **neither complete nor optimal**.

3.2 1

$f(n) = g(n) + h(n)$

- $g(n)$: exact cost from start to n (known, accumulated)
- $h(n)$: estimated cost from n to goal (heuristic)
- $f(n)$: estimated total path cost through n

A* balances exploration of already-cheap paths (g) vs promising paths to goal (h).

3.2.1 1Romania Map: Greedy vs A* Side-by-Side

h=100; [goalnode] (Bu) [below=1.5cm of Fa] Bu
h=0; [-Stealth,dblue] (Ar) – node[left,black] 140 (Si); [-Stealth,dblue] (Ar) – node[right,black] 118 (Ti); [-Stealth,dblue] (Ar) – node[right,black] 75 (Ze); [-Stealth,dblue] (Si) – node[left,black] 99 (Fa); [-Stealth,dblue] (Si) – node[right,black] 80 (Ri); [-Stealth,dblue] (Fa) – node[left,black] 211 (Bu); [-Stealth,dblue] (Ri) – node[right,black] 97 (Pi); [-Stealth,dblue] (Pi) – node[right,black] 101 (Bu); [dred,ultra thick] (Ar) – (Si) – (Fa) – (Bu); [right=0.3cm of Bu,color=dred] **Greedy: cost=450**; [dgreen,ultra thick,dashed] (Ar) to[bend right=8] (Si); [dgreen,ultra thick,dashed] (Si) to[bend right=8] (Ri) – (Pi) to[bend right=8] (Bu); [above=0.3cm of Pi,color=dgreen] **A*:**

cost=418 (optimal!);

A* Trace – Romania (start)

Step	Open List: (node: $g+h=f$)	Expand (lowest f)	Greedy found
1	[(Arad:0+366=366)]	Arad	
2	[(Sibiu:140+253= 393), (Tim:118+329=447), (Zer:75+374=449)]	Sibiu ($f=393$)	
3	[(Rim:220+193= 413), (Fag:239+176=415), (Tim:447), (Zer:449)]	Rimnicu ($f=413$)	
4	[(Fag:415), (Pit:317+100= 417), (Tim:447), (Zer:449)]	Fagaras ($f=415$)	
5	[(Pit: 417), (Buc-Fag:239+211=450), (Tim:447), (Zer:449)]	Pitesti ($f=417$)	
6	[(Buc-Pit:317+101+0= 418), (Buc-Fag:450), ...]	Bucharest $f=418$ = GOAL!	

Arad→Sibiu→Fagaras→Bucharest=450 (3 cities, greedy at each step).
A* found Arad→Sibiu→Rimnicu→Pitesti→Bucharest=418 (4 cities but cheaper). A* is optimal.

3.3 1

$h(n) \leq h^*(n)$ for all n (Admissibility)

$h(n) \leq c(n, a, n') + h(n')$ (Consistency / Triangle Inequality)

Optimality Conditions for A*

Version	Requirement	Why It Works
A* Tree Search	h is admissible	Never overestimates $\Rightarrow$ won't skip optimal path
A* Graph Search	h is consistent (monotone)	$f(n)$ non-decreasing along paths $\Rightarrow$ first dequeue = optimal

Q: [Easy] What is an admissible heuristic? Give an example.

A: $h(n) \leq h^*(n)$: never overestimates the true cost. For 8-puzzle: Manhattan distance is admissible (each tile needs at least that many moves). Euclidean distance is also admissible.

Q: [Medium] 8-puzzle: h_1 =Manhattan distance, h_2 =Euclidean distance. Which is more informed? Both admissible?

A: Both admissible since actual moves $\geq$ any geometric distance. $h_1 \geq h_2$ always (Manhattan $\geq$ Euclidean in grid), so h_1 is more informed. A* with h_1 expands fewer nodes — it “knows more”.

Q: [Hard] Prove that consistency implies A* graph search is optimal.

A: Consistency $\Rightarrow f(n') = g(n') + h(n') \geq g(n) + c(n, a, n') + h(n') \geq g(n) + h(n) = f(n)$. So f is non-decreasing along paths. When A* dequeues node n , all remaining open nodes have $f \geq f(n)$. When goal G is dequeued with cost $g(G)$, since $h(G) = 0$, we have $f(G) = g(G)$. Any other path to G through the open list has $f \geq g(G)$, so no cheaper path exists. **Therefore $g(G)$ is optimal.**

3.4 1

If $h_2(n) \geq h_1(n)$ for all n (and both admissible), then h_2 **dominates** h_1 . A* with h_2 will expand fewer or equal nodes. Always prefer the higher admissible heuristic.

4. 1

Two-player zero-sum games: MAX wins exactly what MIN loses. **MAX** maximizes utility; **MIN** minimizes it. Both play optimally.

4.1 1

$$\text{MINIMAX}(n) = \begin{cases} \text{UTILITY}(n) & n \text{ terminal} \\ \max_{c \in \text{children}(n)} \text{MINIMAX}(c) & n \text{ is MAX node} \\ \min_{c \in \text{children}(n)} \text{MINIMAX}(c) & n \text{ is MIN node} \end{cases}$$

Complexity: Time $O(b^m)$; Space $O(bm)$ (DFS traversal).

4.1.1 1Minimax Full Tree Dry Run

$v = 3$ child node[leaf] (D) 3 child node[leaf] (E) 5 child node[minnode] (C) MIN
 $v = 2$ child node[leaf] (F) 2 child node[leaf] (G) 9 ; [left=0.3cm of B,font=,color=mgray] $\min(3, 5) = 3$; [right=0.3cm of C,font=,color=mgray] $\min(2, 9) = 2$; [right=0.4cm of A,font=,color=dblue] $\max(3, 2) = \mathbf{3}$; [dgreen,very thick,dashed,-Stealth] (A) –
 node[left,font=,color=dgreen] **play** (B);

Minimax Value Table

Node	Type	Children values	Value	Reasoning
D	Terminal	–	3	Leaf utility
E	Terminal	–	5	Leaf utility
F	Terminal	–	2	Leaf utility
G	Terminal	–	9	Leaf utility
B	MIN	D=3, E=5	3	$\min(3, 5) = 3$ — MIN picks worst for MAX
C	MIN	F=2, G=9	2	$\min(2, 9) = 2$
A	MAX	B=3, C=2	3	$\max(3, 2) = 3$ — MAX plays to B

4.2 1

Alpha-Beta eliminates branches that cannot affect the minimax value. Achieves same result as full minimax with fewer node evaluations.

Alpha and Beta Variables

α (**alpha**): Best value MAX can *guarantee* from current path. Initialized to $-\infty$. Updated only at MAX nodes.

β (**beta**): Best value MIN can *guarantee* from current path. Initialized to $+\infty$. Updated only at MIN nodes.

Pruning rules:

At MIN node: if current value $\leq \alpha \Rightarrow$ **prune** (MAX already has better option elsewhere).

At MAX node: if current value $\geq \beta \Rightarrow$ **prune** (MIN already has better option elsewhere).

At MIN node: prune remaining children if $v \leq \alpha$

At MAX node: prune remaining children if $v \geq \beta$

4.2.1 1Alpha-Beta Full Dry Run with Three MAX Children

$v = 3$ child node[minnode] (B) MIN
 $\beta=3$
 $v = 3$ child node[leaf] 3 child node[leaf] 5 child node[minnode] (C) MIN
 $\alpha=3$
 $v = 2$ child node[leaf] 2 child[missing] child node[minnode] (D) MIN

$v = 1$ child node[leaf] 7 child node[leaf] 1 ; [draw=dred,dashed,fill=pink,minimum width=9mm,minimum height=7mm,font=] (pruned) at (7.5,-4.2) 9; [dred,dashed,thick] (C) – (pruned); [right=0.1cm of pruned,font=,color=dred]


PRUNED ($2 \leq \alpha=3$); [dgreen,very thick,dashed,-Stealth] (A) – node[left,font=,color=dgreen] **play to B** (B);

Alpha-Beta Detailed Trace

Step	Node	Action	α	β	Result
1	A(MAX)	Initialize, visit B	$-\infty$	$+\infty$	Call B
2	B(MIN)	Leaf 3: $v = 3$, $\beta = \min(+\infty, 3) = 3$	$-\infty$	3	Partial B=3
3	B(MIN)	Leaf 5: $\min(3, 5) = 3$. B done	$-\infty$	3	B=3
4	A(MAX)	B=3: $\alpha = \max(-\infty, 3) = 3$. Visit C	3	$+\infty$	$\alpha = 3$
5	C(MIN)	Leaf 2: $v = 2$. Is $v \leq \alpha$? $2 \leq 3$?	3	$+\infty$	YES! Prune child 9. C=2
6	A(MAX)	C=2: $\max(3, 2) = 3$. Visit D	3	$+\infty$	α stays 3
7	D(MIN)	Leaf 7: $\beta = \min(+\infty, 7) = 7$	3	7	D partial=7
8	D(MIN)	Leaf 1: $\min(7, 1) = 1$. D=1	3	7	D=1
9	A(MAX)	$\max(3, 2, 1) = 3$. Done	3	$+\infty$	Final=3. Play to B.

4.2.2 1Alpha-Beta Best/Worst/Average Case Graph

Alpha-Beta vs Minimax: Nodes Examined ($b = 3$)



Best case (successors sorted best-first): only $O(b^{m/2})$ nodes – doubles effective search depth within same budget. **Worst case** (sorted worst-first): $O(b^m)$ – no pruning at all. **Move ordering** is critical for alpha-beta effectiveness.

Q: [Easy] What do α and β represent?

A: α =best guaranteed value for MAX; β =best guaranteed value for MIN. Prune at MIN when $v \leq \alpha$; at MAX when $v \geq \beta$.

Q: [Medium] Full minimax examines 10^9 nodes. With optimal alpha-beta?

A: Best case: $O(b^{m/2}) = (b^m)^{0.5} = (10^9)^{0.5} \approx 31,623$ nodes. About **31,600×** reduction.

Q: [Hard] A(MAX)→B(MIN),C(MIN). B has leaves 6,4,9; C has leaves 2,8,3. Trace alpha-beta with $\alpha = -\infty$, $\beta = +\infty$. Which nodes are pruned?

A: B: leaf 6→ $\beta = 6$; leaf 4→ $\beta = 4$; leaf 9: $\min(4, 9) = 4$. B=4. Back to A: $\alpha = \max(-\infty, 4) = 4$. C: leaf 2→ $v = 2$. Check $v \leq \alpha$? $2 \leq 4$? **YES** ⇒ **prune leaves 8 and 3**. C=2. A= $\max(4, 2) = 4$. Play to B.
Pruned: nodes 8 and 3 under C.

5. 1


Local search works on a **single current state** with no frontier or path memory. It moves to neighbors to improve an **objective function** $f(n)$.

Advantages: Uses very little memory $O(1)$; works for continuous/very large spaces.

Disadvantages: Can get stuck at local optima, plateaus, or ridges.

5.0.1 1Objective Function Landscape – Problems Illustrated

Local Search Landscape: Local Optima, Plateaus, Ridges, Global Optimum




Hill Climbing Variants Compared

Variant	How It Works	Pros / Cons
Simple HC	Move to FIRST improving neighbor	Fast per step; may miss better neighbors
Steepest Ascent HC	Evaluate ALL neighbors; move to BEST	Better quality; still stuck at local optima
Stochastic HC	Pick randomly among uphill moves	Less trapping; slower convergence
Random Restart HC	Repeat from random start states	Probabilistically complete
Simulated Annealing	Accept downhill with prob $e^{\Delta E/T}$, $T \rightarrow 0$	Escapes local optima; guaranteed global opt

5.0.2 1Steepest Ascent HC Dry Run with Graph

Objective Values by State: Steepest Ascent Trace



Steepest Ascent HC: Start

Step	Current	f	Neighbors	f(neighbors)	Move to
1	A	2	B	f(B)=5 > 2	B
2	B	5	A, C	f(A)=2, f(C)=9	C (best)
3	C	9	B, D	f(B)=5, f(D)=4	Both < 9: LOCAL (=GLOBAL) OPTI- MUM

Path: A→B→C.

Found global optimum in 2 steps. (Convenient here; usually local optima block progress.)

Simulated Annealing: Escape Local Optima

Accept worse states with probability $P = e^{\Delta E/T}$ where $\Delta E = f(\text{new}) - f(\text{current}) < 0$ and T is temperature (decreasing).

When T is high (early): $P \approx 1$ – accept almost any move (explore).

When $T \rightarrow 0$ (later): $P \rightarrow 0$ – only accept improvements (exploit).

Cooling schedule: $T(t) = T_0 \cdot \alpha^t$ (geometric) or $T_0 / \log(t + 1)$ (logarithmic). Logarithmic cooling guarantees convergence to global optimum.

Q: [Easy] What is a local optimum? Why does HC get stuck?

A: A local optimum is a state where all neighbors have lower f value but it is not the global maximum. HC gets stuck because it only moves to neighbors with higher f , so it cannot escape.

Q: [Medium] States: $A(4) \leftrightarrow B(7) \leftrightarrow C(5) \leftrightarrow D(9) \leftrightarrow E(3)$. Start=A. Steepest ascent HC. Does it find global?

A: $A(4) \rightarrow B(7)$. At B: neighbors $A(4) < 7$ and $C(5) < 7$. **HC stops at B(7)**. D(9) is the global optimum but HC can't reach it without going down through C(5). HC fails here.


Q: [Hard] Why is random-restart HC probabilistically complete? If $p = 0.2$, expected restarts?

A: Each restart independently finds the global optimum with probability $p > 0$. Probability of never succeeding in k restarts $= (1 - p)^k \rightarrow 0$. Expected restarts $= 1/p = 1/0.2 = 5$. It is probabilistically complete (probability of finding optimum approaches 1 as restarts increase).

6.1

GAs maintain a **population** of candidate solutions. Evolution via **selection**, **crossover**, and **mutation** over many generations.

6.0.1 1GA Cycle Diagram



6.1 1


6.1.1 1Roulette Wheel (Fitness Proportionate) Selection

$$P(i) = \frac{\text{fitness}(i)}{\sum_{j=1}^N \text{fitness}(j)}$$

Roulette Wheel Dry Run

Population: A(4), B(2), C(6), D(8). Total fitness=20.

$P(A) = 4/20 = 0.20$, $P(B) = 2/20 = 0.10$, $P(C) = 6/20 = 0.30$, $P(D) = 8/20 = 0.40$



Cumulative ranges:

A: [0.00, 0.20)

B: [0.20, 0.30)

C: [0.30, 0.60)

D: [0.60, 1.00)

$r = 0.45 \rightarrow \mathbf{C}$ selected

$r = 0.72 \rightarrow \mathbf{D}$ selected

$r = 0.18 \rightarrow \mathbf{A}$ selected

6.1.2 1Tournament Selection

Tournament Selection (k)

Pick k individuals at random; the one with highest fitness wins.

Tournament 1: randomly pick C(6) and A(4). Winner: **C(6)**.

Tournament 2: randomly pick D(8) and B(2). Winner: **D(8)**.

Parents for crossover: C and D. Higher $k \Rightarrow$ more selection pressure.


6.1.3 1Rank-Based Selection

$$P(\text{rank } r) = \frac{r}{N(N+1)/2}$$


Sort individuals by fitness; assign rank 1 (worst) to N (best). Avoids domination by super-fit individuals.

6.2 1

6.2.1 1Single-Point Crossover




6.2.2 1Bit-Flip Mutation



Mutation probability p_m (typically 0.001 to 0.01) applied to each bit independently.

6.2.3 1GA Fitness Over Generations Graph

Typical GA Convergence: Best/Average Fitness Over Generations



Q: [Medium] Population A(10), B(20), C(30), D(40). Roulette probabilities and cumulative ranges?

A: Total=100. $P(A)=10\%$: [0.00,0.10). $P(B)=20\%$: [0.10,0.30). $P(C)=30\%$: [0.30,0.60). $P(D)=40\%$: [0.60,1.00). $r = 0.55 \rightarrow C$. $r = 0.85 \rightarrow D$.

Q: [Hard] P1=110010, P2=001101. 2-point crossover after positions 2 and 4. Then bit-flip O1 at position 5.

A: Segments: P1=[11—00—10], P2=[00—11—01]. Swap middle segment: O1=11+11+10=**111110**. O2=00+00+01=**000001**. Mutation O1 position 5: flip 1→0: **111100**.

7. 1

Types of Machine Learning

Type	Description	Examples
Supervised	Labeled data: learn $f : X \rightarrow Y$	Classification, Regression
Unsupervised	Find structure in unlabeled data	Clustering, Dim. Reduction
Reinforcement	Agent learns from environment rewards	Game playing, Robotics


$$y = w_1x + w_0 \quad \text{(Simple Linear Regression)}$$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$w_1 = \frac{n \sum x_i y_i - \sum x_i \sum y_i}{n \sum x_i^2 - (\sum x_i)^2}, \quad w_0 = \bar{y} - w_1 \bar{x}$$

7.0.1 1Linear Regression Dry Run with Scatter Plot

Scatter Plot and Regression Line




Computation Table: x

	x_i	y_i	x_i^2	$x_i y_i$	
$y=[2,4,5,4,5]$	1	2	1	2	$w_1 = \frac{5(66) - 15(20)}{5(55) - 15^2} = \frac{330 - 300}{275 - 225} = \frac{30}{50} = 0.6$
	2	4	4	8	
	3	5	9	15	
	4	4	16	16	
	5	5	25	25	
	15	20	55	66	
					$\bar{x} = 15/5 = 3, \bar{y} = 20/5 = 4$
					$w_0 = 4 - 0.6(3) = 4 - 1.8 = 2.2$
					$\hat{y} = 0.6x + 2.2$
					Check: $\hat{y}(1) = 2.8, \hat{y}(3) = 4.0, \hat{y}(5) = 5.2$

8. 1

8.1 1




Key facts: Each connection has a weight. Each neuron (except input) has a bias. Learning = adjusting weights and biases via backpropagation. More hidden layers = deeper network = more complex features.

8.2 1

8.2.1 1All Major Activation Functions

Activation Functions: Sigmoid, Tanh, ReLU, Linear



Sigmoid in Detail – Used in Backpropagation

$$\sigma(z) = \frac{1}{1 + e^{-z}}, \quad \sigma'(z) = \sigma(z) \cdot (1 - \sigma(z)) \leq 0.25$$

Key values: $\sigma(0) = 0.5$, $\sigma(1) \approx 0.731$, $\sigma(-1) \approx 0.269$, $\sigma(2) \approx 0.880$.

Maximum derivative: $\sigma'(0) = 0.5 \times 0.5 = 0.25$ (at $z = 0$).

Vanishing gradient: for $|z| > 3$, $\sigma' \approx 0$, gradients shrink and deep layers learn very slowly.

Activation Function Comparison Table

Function	Formula	Range	Properties & Issues
Sigmoid	$1/(1 + e^{-z})$	$(0, 1)$	Smooth; probabilistic output; saturates at extremes; max $\sigma' \leq 0.25$; vanishing gradient in deep nets
Tanh	$(e^z - e^{-z})/(e^z + e^{-z})$	$(-1, 1)$	Zero-centered (better gradient flow than sigmoid); still saturates
ReLU	$\max(0, z)$	$[0, \infty)$	No saturation for $z > 0$; fast to compute; dying ReLU: neurons with $z < 0$ always output 0
Leaky ReLU	z if $z > 0$, $0.01z$ otherwise	$(-\infty, \infty)$	Fixes dying ReLU; small gradient for $z < 0$
Softmax	$e^{z_k} / \sum_j e^{z_j}$	$(0, 1)$	Multi-class output layer; outputs sum to 1 (probability distribution)


8.3 1

A single perceptron computes $\hat{y} = \text{step}(\mathbf{w} \cdot \mathbf{x} + b)$. The decision boundary is a hyperplane $\mathbf{w} \cdot \mathbf{x} + b = 0$.

$$\hat{y} = \text{step}\left(\sum_i w_i x_i + b\right), \quad \text{step}(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$$

$$w_i \leftarrow w_i + \eta(y - \hat{y})x_i, \quad b \leftarrow b + \eta(y - \hat{y}) \quad (\text{Perceptron Update})$$

XOR is NOT Linearly Separable (no line can separate ○ from ●)



AND and OR are linearly separable (single perceptron works). XOR requires at least one hidden layer. This was a major result in AI history (Minsky & Papert, 1969) that temporarily halted ANN research.

9. 1

Backpropagation applies the **chain rule** of calculus to compute gradients of the loss with respect to every weight. We then do **gradient descent** to minimize loss.

Core Notation

$z_j^{(l)}$: pre-activation at neuron j , layer l ($z = \sum_i w_{ij}a_i + b_j$)

$a_j^{(l)} = \sigma(z_j^{(l)})$: post-activation (output of neuron)

$\delta_j^{(l)} = \partial L / \partial z_j^{(l)}$: local gradient / error signal (delta)

Loss: $L = \frac{1}{2}(y - a_{\text{out}})^2$ (MSE, factor $\frac{1}{2}$ for clean derivative)

Update: $w \leftarrow w - \eta \cdot \frac{\partial L}{\partial w}$ (η = learning rate)

$$\delta_{\text{out}} = (a_{\text{out}} - y) \cdot \sigma'(z_{\text{out}}) \quad (\text{Output delta})$$

$$\delta_j^{(l)} = \sigma'(z_j^{(l)}) \cdot \sum_k \delta_k^{(l+1)} \cdot w_{jk} \quad (\text{Hidden delta, chain rule})$$

$$\frac{\partial L}{\partial w_{ij}} = \delta_j^{(l)} \cdot a_i^{(l-1)} \quad (\text{Gradient for weight } w_{ij})$$

9.1 1

9.1.1 1Network Architecture Diagram

Step 1: Forward Pass

Hidden layer pre-activations:

$$z_{h1} = w_{11}x_1 + w_{21}x_2 + b_h = 0.15(0.05) + 0.25(0.10) + 0.35 = 0.0075 + 0.025 + 0.35 = \mathbf{0.3825}$$

$$z_{h2} = w_{12}x_1 + w_{22}x_2 + b_h = 0.20(0.05) + 0.30(0.10) + 0.35 = 0.010 + 0.030 + 0.35 = \mathbf{0.3900}$$

Hidden layer activations (Sigmoid):

$$a_{h1} = \sigma(0.3825) = \frac{1}{1 + e^{-0.3825}} = \frac{1}{1 + 0.6820} = \frac{1}{1.6820} = \mathbf{0.5944}$$

$$a_{h2} = \sigma(0.3900) = \frac{1}{1 + e^{-0.3900}} = \frac{1}{1 + 0.6771} = \frac{1}{1.6771} = \mathbf{0.5963}$$

Output layer:

$$\begin{aligned} z_o &= w_{1o} \cdot a_{h1} + w_{2o} \cdot a_{h2} + b_o = 0.40(0.5944) + 0.45(0.5963) + 0.60 \\ &= 0.23776 + 0.26834 + 0.60 = \mathbf{1.1061} \end{aligned}$$

$$a_o = \sigma(1.1061) = \frac{1}{1 + e^{-1.1061}} = \frac{1}{1 + 0.3307} = \frac{1}{1.3307} = \mathbf{0.7514}$$

$$\text{Loss: } L = \frac{1}{2}(y - a_o)^2 = \frac{1}{2}(0.01 - 0.7514)^2 = \frac{1}{2}(-0.7414)^2 = \frac{1}{2}(0.5497) = \mathbf{0.2748}$$

Step 2: Backward Pass – Output Layer**Sigmoid derivative at output:**

$$\sigma'(z_o) = a_o(1 - a_o) = 0.7514 \times 0.2486 = \mathbf{0.1868}$$

Output delta (chain rule: $\partial L / \partial z_o = \partial L / \partial a_o \cdot \partial a_o / \partial z_o$):

$$\frac{\partial L}{\partial a_o} = a_o - y = 0.7514 - 0.01 = 0.7414$$

$$\delta_o = \frac{\partial L}{\partial a_o} \cdot \sigma'(z_o) = 0.7414 \times 0.1868 = \mathbf{0.1385}$$

Gradients for output weights:

$$\frac{\partial L}{\partial w_{1o}} = \delta_o \cdot a_{h1} = 0.1385 \times 0.5944 = \mathbf{0.08232}$$

$$\frac{\partial L}{\partial w_{2o}} = \delta_o \cdot a_{h2} = 0.1385 \times 0.5963 = \mathbf{0.08258}$$

$$\frac{\partial L}{\partial b_o} = \delta_o \cdot 1 = \mathbf{0.1385}$$

Step 3: Backward Pass – Hidden Layer**Propagate delta back through output weights:**

$$\frac{\partial L}{\partial a_{h1}} = \delta_o \cdot w_{1o} = 0.1385 \times 0.40 = 0.05540$$

$$\frac{\partial L}{\partial a_{h2}} = \delta_o \cdot w_{2o} = 0.1385 \times 0.45 = 0.06233$$

Sigmoid derivatives at hidden layer:

$$\sigma'(z_{h1}) = a_{h1}(1 - a_{h1}) = 0.5944 \times 0.4056 = \mathbf{0.2410}$$

$$\sigma'(z_{h2}) = a_{h2}(1 - a_{h2}) = 0.5963 \times 0.4037 = \mathbf{0.2407}$$

Hidden deltas:

$$\delta_{h1} = \frac{\partial L}{\partial a_{h1}} \cdot \sigma'(z_{h1}) = 0.05540 \times 0.2410 = \mathbf{0.013351}$$

$$\delta_{h2} = \frac{\partial L}{\partial a_{h2}} \cdot \sigma'(z_{h2}) = 0.06233 \times 0.2407 = \mathbf{0.015015}$$

Gradients for hidden weights:

$$\frac{\partial L}{\partial w_{11}} = \delta_{h1} \cdot x_1 = 0.013351 \times 0.05 = \mathbf{0.000668}$$

$$\frac{\partial L}{\partial w_{21}} = \delta_{h1} \cdot x_2 = 0.013351 \times 0.10 = \mathbf{0.001335}$$

$$\frac{\partial L}{\partial w_{12}} = \delta_{h2} \cdot x_1 = 0.015015 \times 0.05 = \mathbf{0.000751}$$

$$\frac{\partial L}{\partial w_{22}} = \delta_{h2} \cdot x_2 = 0.015015 \times 0.10 = \mathbf{0.001502}$$

Step 4: Weight Updates (η)

$$w_{1o} \leftarrow 0.40 - 0.5(0.08232) = 0.40 - 0.04116 = \mathbf{0.35884}$$

$$w_{2o} \leftarrow 0.45 - 0.5(0.08258) = 0.45 - 0.04129 = \mathbf{0.40871}$$

$$b_o \leftarrow 0.60 - 0.5(0.1385) = 0.60 - 0.06925 = \mathbf{0.53075}$$


$$w_{11} \leftarrow 0.15 - 0.5(0.000668) = 0.15 - 0.000334 = \mathbf{0.149666}$$

$$w_{21} \leftarrow 0.25 - 0.5(0.001335) = 0.25 - 0.000668 = \mathbf{0.249332}$$

$$w_{12} \leftarrow 0.20 - 0.5(0.000751) = 0.20 - 0.000376 = \mathbf{0.199624}$$

$$w_{22} \leftarrow 0.30 - 0.5(0.001502) = 0.30 - 0.000751 = \mathbf{0.299249}$$

9.1.2 Loss Reduction Over Training Epochs



After the one weight update above, loss drops from 0.2748 to ≈ 0.2468 . Training repeats this cycle for many epochs (typically hundreds to thousands) until convergence.

9.2 1

$$\delta_j^{(l)} = \sigma'(z_j^{(l)}) \cdot \sum_k \delta_k^{(l+1)} \cdot w_{jk}^{(l+1)} \quad (\text{General delta - chain rule})$$

$$\frac{\partial L}{\partial w_{ij}^{(l)}} = \delta_j^{(l)} \cdot a_i^{(l-1)} \quad (\text{Gradient for any weight})$$

The key insight: delta at layer l depends on deltas at layer $l+1$ (already computed), multiplied by the connecting weights and the local sigmoid derivative. This is why we go **backward** from output to input.

Q: [Easy] What is $\sigma'(z)$ for sigmoid? If $\sigma(z) = 0.7$, compute $\sigma'(z)$.

A: $\sigma'(z) = \sigma(z)(1 - \sigma(z))$. If $\sigma(z) = 0.7$: $\sigma' = 0.7 \times 0.3 = 0.21$.

Q: [Medium] Why multiply by $\sigma'(z)$ when computing δ ?

A: By the chain rule: $\delta = \frac{\partial L}{\partial z} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial z} = \frac{\partial L}{\partial a} \cdot \sigma'(z)$. The $\sigma'(z)$ term measures how much the pre-activation z affects the output a . Without it, we'd have the wrong gradient.

Q: [Hard] Network: $x = 0.6$, 1 hidden ($w_h = 0.5$, $b_h = 0$), 1 output ($w_o = 0.8$, $b_o = 0$). $y = 1$, $\eta = 0.1$. Full forward and backward pass.


A: Forward: $z_h = 0.5(0.6) = 0.3$, $a_h = \sigma(0.3) = 0.5744$. $z_o = 0.8(0.5744) = 0.4595$, $a_o = \sigma(0.4595) = 0.6129$. $L = \frac{1}{2}(1 - 0.6129)^2 = 0.0750$.

Backward: $\sigma'(z_o) = 0.6129(0.3871) = 0.2373$. $\delta_o = (0.6129 - 1)(0.2373) = (-0.3871)(0.2373) = -0.09189$. $\partial L / \partial w_o = \delta_o \cdot a_h = -0.09189(0.5744) = -0.05279$. $w_o \leftarrow 0.8 - 0.1(-0.05279) = \mathbf{0.80528}$.

$\sigma'(z_h) = 0.5744(0.4256) = 0.2445$. $\delta_h = (\delta_o \cdot w_o) \cdot \sigma'(z_h) = (-0.09189 \times 0.8)(0.2445) = -0.07351 \times 0.2445 = -0.017972$. $\partial L / \partial w_h = \delta_h \cdot x = -0.017972(0.6) = -0.010783$. $w_h \leftarrow 0.5 - 0.1(-0.010783) = \mathbf{0.501078}$.

10. 1


10.1 1



Underfitting vs Good Fit vs Overfitting

Condition	Description	Train Err	Test Err	Solutions
Underfitting	Model too simple; misses patterns	High	High	More complexity, more features, fewer constraints
Good Fit	Captures true signal	Low	Low	–
Overfitting	Memorizes noise; no generalization	Very Low	High	Regularization (L_1/L_2), Dropout, Early stopping, More training data



10.2 1



Each sigmoid layer multiplies gradients by at most $\sigma' \leq 0.25$. After 10 layers: $0.25^{10} \approx 10^{-6}$. Early layers receive almost zero gradient and barely learn. This is why deep networks use ReLU instead of sigmoid in hidden layers.

10.3 1

Without Dropout

With Dropout ($p = 0.5$)

11. 1

11.1 1

		Predicted	
		Positive	Negative
Actual	Positive	TP True Positive	FN False Negative (Type II)
	Negative	FP False Positive (Type I)	TN True Negative

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP} \quad \text{"Of all predicted positive, how many actually are?"}$$

$$\text{Recall (Sensitivity)} = \frac{TP}{TP + FN} \quad \text{"Of all actual positive, how many did we catch?"}$$

$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}$$


11.1.1 1Numeric Dry Run

Example: 10 samples (5 positive)

		Predicted	
		Positive	Negative
Actual	Positive	TP=3	FN=2
	Negative	FP=2	TN=3

$$\text{Accuracy} = \frac{3+3}{10} = 0.60 \quad \text{Precision} = \frac{3}{3+2} = 0.60 \quad \text{Recall} = \frac{3}{3+2} = 0.60 \quad F_1 = \frac{2(0.6)(0.6)}{0.6+0.6} = 0.60$$

11.1.2 1Precision vs Recall Trade-off



11.2 1

3-Class Confusion Matrix (Cat / Dog / Bird)

Actual \ Pred	Cat	Dog	Bird	Class	TP	FP	FN	Prec
Cat (10)	8	1	1	Cat	8	2	2	0.800
Dog (10)	2	7	1	Dog	7	3	3	0.700
Bird (10)	0	2	8	Bird	8	2	2	0.800

Macro Average: Average per-class metrics (treats each class equally).

Macro Prec = $(0.800 + 0.700 + 0.800)/3 = 0.767$. Macro Rec = 0.767 . Macro $F_1 = 0.767$.

Micro Average: Aggregate all TP/FP/FN across classes (treats each sample equally).

$TP_{\text{tot}} = 8 + 7 + 8 = 23$, $FP_{\text{tot}} = 2 + 3 + 2 = 7$, $FN_{\text{tot}} = 2 + 3 + 2 = 7$.

Micro Prec = $23/30 = 0.767$. Micro Rec = $23/30 = 0.767$. Micro $F_1 = 0.767$.

Macro treats all classes equally; use when classes are equally important. Micro is dominated by larger classes.

Q: [Easy] What does FN mean in cancer screening?

A: Patient HAS cancer but model says they DON'T. A False Negative (Type II error / missed diagnosis) is the most dangerous error in medical screening.

Q: [Medium] When prefer Recall over Precision?

A: When missing a true positive (FN) is more costly than a false alarm (FP). Examples: cancer screening (missing cancer is catastrophic), fraud detection (missing fraud is costly). Prefer Precision when false alarms are costly (e.g., spam filtering — don't want valid emails flagged).

Q: [Hard] 4-class matrix rows: $C1=[50,2,3,0]$, $C2=[1,40,4,2]$, $C3=[2,3,45,1]$, $C4=[0,1,2,48]$. Compute Macro F_1 . Which class is worst?

A: $TP=[50,40,45,48]$. $FP(\text{col sum} - TP)$: $C1: 3$, $C2: 6$, $C3: 9$, $C4: 3$. $FN(\text{row sum} - TP)$: $C1: 5$, $C2: 7$, $C3: 6$, $C4: 3$. P : $50/53=0.943$, $40/46=0.870$, $45/54=0.833$, $48/51=0.941$. R : $50/55=0.909$, $40/47=0.851$, $45/51=0.882$, $48/51=0.941$. F_1 : 0.926 , 0.860 , 0.857 , 0.941 . Macro $F_1 = (0.926 + 0.860 + 0.857 + 0.941)/4 = 0.896$. **Poorest: C3** ($F_1 = 0.857$).

12. 1


$$\text{Output size} = \left\lfloor \frac{N - F + 2P}{S} \right\rfloor + 1$$

N =input size, F =filter/kernel size, P =zero-padding, S =stride.

$$\text{Parameters per conv layer} = (F \times F \times C_{\text{in}} + 1) \times \text{num_filters}$$

The +1 is for the bias of each filter.

12.1 1




Computing output[0][0] (top-left 3x3 patch of input convolved with filter):


$$\text{out}[0][0] = 1(1) + 2(0) + 3(-1) + 5(1) + 4(0) + 3(-1) + 1(1) + 2(0) + 3(-1) = (1-3) + (5-3) + (1-3) = -2 + 2 - 2 = -2$$

Output size: $\lfloor (5 - 3 + 0) / 1 \rfloor + 1 = 3$. So 3x3 output from 5x5 input with 3x3 filter, $P = 0$, $S = 1$.

12.2 1



12.2.1 1 Full CNN Architecture Sizes and Parameters



CNN Architecture: Input $32 \times 32 \times 3$

Layer	Output Size	Formula	Parameters
Input	$32 \times 32 \times 3$	—	0
Conv1: 16f, 5×5 , $P=2$, $S=1$	$32 \times 32 \times 16$	$(32 - 5 + 4)/1 + 1 = 32$	$(5 \times 5 \times 3 + 1) \times 16 = 1216$
MaxPool1: 2×2 , $S=2$	$16 \times 16 \times 16$	$32/2 = 16$	0
Conv2: 32f, 3×3 , $P=1$, $S=1$	$16 \times 16 \times 32$	$(16 - 3 + 2)/1 + 1 = 16$	$(3 \times 3 \times 16 + 1) \times 32 = 4640$
MaxPool2: 2×2 , $S=2$	$8 \times 8 \times 32$	$16/2 = 8$	0
Flatten	2048	$8 \times 8 \times 32 = 2048$	0
Total params	—	—	5856

Q: [Easy] Input 32×32 , filter 5×5 , $P=0$, $S=1$. Output size?

A: $(32 - 5 + 0)/1 + 1 = 28$. Output is 28×28 .

Q: [Medium] Why does max pooling give translation invariance?

A: The max over a region is unchanged by small shifts of the feature within that region. An edge detector firing slightly left or right gives the same pooled output — the network doesn't care exactly where the feature is, just that it's present.

Q: [Hard] Input $64 \times 64 \times 3$. Conv: 32 filters, 3×3 , $P=1$, $S=1$. Pool: 2×2 , $S=2$. Conv: 64 filters, 3×3 , $P=1$, $S=1$. FC: 512 neurons. FC: 10 (output). Compute all sizes and total parameters.

A: Conv1: $(64 - 3 + 2)/1 + 1 = 64$. Out= $64 \times 64 \times 32$. Params= $(3 \times 3 \times 3 + 1) \times 32 = 896$. Pool: $32 \times 32 \times 32$. Conv2: $(32 - 3 + 2)/1 + 1 = 32$. Out= $32 \times 32 \times 64$. Params= $(3 \times 3 \times 32 + 1) \times 64 = 18496$. Pool2: $16 \times 16 \times 64$. Flatten= 16384 . FC1: $16384 \times 512 + 512 = 8,389,120$. FC2: $512 \times 10 + 10 = 5130$. **Total: 8,413,642 parameters.**


13. 1

13.1 1

K-Means Algorithm

1. Initialize k centroids randomly (or K-Means++ initialization)
2. **Assignment step:** Assign each point to nearest centroid ($\arg \min_c \|x - \mu_c\|^2$)
3. **Update step:** Recompute centroids as mean of assigned points
4. Repeat steps 2–3 until assignments stop changing (convergence)

Objective: Minimize within-cluster sum of squares (WCSS) = $\sum_c \sum_{x \in c} \|x - \mu_c\|^2$.

K-Means: $k = 2$, 2 Iterations to Convergence

K-Means Full Trace: k

Initial centroids: $C1=(1,1)$, $C2=(5,4)$.

Iteration 1 – **Assignment** (Euclidean distance $d = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$):

Point	Dist to $C1=(1,1)$	Dist to $C2=(5,4)$	Assign
A(1,1)	0.00	5.00	C1
B(1,2)	1.00	4.47	C1
C(2,1)	1.00	4.24	C1
D(5,4)	5.00	0.00	C2
E(5,5)	5.66	1.00	C2
F(6,4)	5.83	1.00	C2

Update centroids: $C1 = (\frac{1+1+2}{3}, \frac{1+2+1}{3}) = (\frac{4}{3}, \frac{4}{3}) = (1.33, 1.33)$ $C2 = (\frac{5+5+6}{3}, \frac{4+5+4}{3}) = (\frac{16}{3}, \frac{13}{3}) = (5.33, 4.33)$

Iteration 2: Re-assign with new centroids. All points remain in same clusters $\Rightarrow$ **CONVERGED!**

Final clusters: $\{A, B, C\}$ and $\{D, E, F\}$.


WCSS $= d(A, C1)^2 + d(B, C1)^2 + d(C, C1)^2 + d(D, C2)^2 + d(E, C2)^2 + d(F, C2)^2 = (0.22 + 0.56 + 0.56 + 0.22 + 0.56 + 0.44) = 2.56$

13.2 1**Linkage Criteria**

Linkage	Formula	Behavior
Single (MIN)	$\min_{a \in A, b \in B} d(a, b)$	Chaining effect; elongated clusters
Complete (MAX)	$\max_{a \in A, b \in B} d(a, b)$	Compact, roughly equal-sized clusters
Average (UPGMA)	$\frac{1}{ A B } \sum_{a,b} d(a, b)$	Balance between single and complete
Ward's	Minimizes WCSS increase	Best overall; minimizes variance

13.2.1 1Complete-Linkage Agglomerative Dry Run**5 Points: A**

			A	B	C	D	E	
Initial	distance	matrix:	A	0	1	3	7	9
			B	1	0	2	6	8
			C	3	2	0	4	6
			D	7	6	4	0	2
			E	9	8	6	2	0
Step	Action (merge min distance pair)		Clusters		Height			
1	Merge A+B (d=1, minimum)		{AB}, {C}, {D}, {E}		1			
2	Merge D+E (d=2)		{AB}, {C}, {DE}		2			
3	Merge AB+C: complete linkage $\max(d(A, C), d(B, C)) = \max(3, 2) = 3$		{ABC}, {DE}		3			
4	Merge ABC+DE: $\max(d(A, D), d(A, E), d(B, D), d(B, E), d(C, D), d(C, E)) = \max(7, 9, 6, 8, 4, 6) = 9$		{ABCDE}		9			



Q: [Easy] What are the two steps of K-Means per iteration?

A: **Assignment:** assign each point to nearest centroid. **Update:** recompute centroids as mean of assigned points.

Q: [Medium] Points P1(2,3), P2(5,4), P3(1,2), P4(6,5). $k = 2$, $C1=(2,3)$, $C2=(6,5)$. One iteration.

A: P1: to $C1=0$, to $C2=\sqrt{25} = 5 \rightarrow C1$. P2: to $C1=\sqrt{10} = 3.16$, to $C2=\sqrt{2} = 1.41 \rightarrow C2$. P3: to $C1=\sqrt{2} = 1.41$, to $C2=\sqrt{34} = 5.83 \rightarrow C1$. P4: to $C1=\sqrt{25} = 5$, to $C2=0 \rightarrow C2$. New $C1 = ((2+1)/2, (3+2)/2) = (1.5, 2.5)$. New $C2 = ((5+6)/2, (4+5)/2) = (5.5, 4.5)$.

Q: [Hard] 1D points $A=1$, $B=2$, $C=4$, $D=8$, $E=9$. Complete-linkage agglomerative clustering. Where to cut for 3 clusters?

A: $d(A,B)=1$, $d(D,E)=1$ (tie; merge both). Step1: Merge $A+B$, $h=1$. Step2: Merge $D+E$, $h=1$. Clusters: $\{AB\}, \{C\}, \{DE\}$. $d(AB,C)=\max(3,2)=3$. $d(C,DE)=\max(5,6)=6$ (using $C=4$, $D=8$, $E=9$). $d(AB,DE)=\max(7,8,6,7)=8$. Step3: $\min=3$, merge $AB+C$, $h=3$. Step4: Merge $ABC+DE$, $h=\max(7,8,5,6)=8$. For 3 clusters: cut between $h=1$ and $h=3 \rightarrow \{AB\}, \{C\}, \{DE\}$.

13. 1

Search Algorithms

A*: $f(n) = g(n) + h(n)$. Admissible: $h(n) \leq h^*(n)$. Consistent: $h(n) \leq c(n, a, n') + h(n')$.

DFS Space: $O(bm)$. BFS/IDS Time: $O(b^d)$. IDS Space: $O(bd)$. UCS optimal for weighted graphs.

Neural Networks & Backpropagation

$$\sigma(z) = \frac{1}{1+e^{-z}}. \sigma'(z) = \sigma(z)(1 - \sigma(z)) \leq 0.25.$$

$$\delta_{\text{out}} = (a_{\text{out}} - y)\sigma'(z_{\text{out}}). \delta_h = (\sum_k \delta_k w_{hk})\sigma'(z_h). w \leftarrow w - \eta \delta a_{\text{prev}}.$$

$$\text{Loss: } L = \frac{1}{2}(y - \hat{y})^2. \text{ Update: } w \leftarrow w - \eta \frac{\partial L}{\partial w}.$$

Performance Metrics

$$\text{Accuracy} = \frac{TP+TN}{N}. \text{ Precision} = \frac{TP}{TP+FP}. \text{ Recall} = \frac{TP}{TP+FN}. F_1 = \frac{2PR}{P+R}.$$

Macro F1: average per-class F1. Micro F1: pool all TP/FP/FN first.

CNN

Output size: $\lfloor (N - F + 2P)/S \rfloor + 1$. Params: $(F^2 \cdot C_{\text{in}} + 1) \times \text{filters}$.

Linear Regression

$$w_1 = \frac{n \sum x_i y_i - \sum x_i \sum y_i}{n \sum x_i^2 - (\sum x_i)^2}. w_0 = \bar{y} - w_1 \bar{x}.$$

Genetic Algorithm

Roulette: $P(i) = \frac{f_i}{\sum f_j}$. Rank: $P(r) = \frac{r}{N(N+1)/2}$. Simulated Annealing: $P(\text{accept worse}) = e^{\Delta E/T}$.